
EL2805 Lab2 Report

jiaying yang
jiayingy@kth.se

940816-0258
yipingx@kth.se

1 Problem 1: Deep Reinforcement Learning for Cartpole

1.1 a) Formulate the RL problem

State Space:

$$S = (P, V, \theta, V_{tip}), -2.4m \leq P \leq 2.4m, V, V_{tip} \in R, -41.8^\circ \leq \theta \leq 41.8^\circ$$

where P is cart position, V is cart velocity, θ is pole angle and V_{tip} is pole velocity at tip.

Action Space:

$$A = \{0, 1\}$$

where 0 indicates pushing cart to the left with 1N, and 1 indicates pushing cart to the right with 1N.

Termination:

When the episode number is 200 or the cart position or pole angle is out of the range.

Rewards:

$$R(-12.5^\circ \leq \theta \leq 12.5^\circ, a = \cdot) = 1$$

$$R(\theta \leq -12.5^\circ, a = \cdot) = 0$$

$$R(\theta \geq 12.5^\circ, a = \cdot) = 0$$

Q-learning works perfectly fine if we have a limited state space and action space. However, when the state space become much bigger, or even there are unlimited number of states like our problem, we need to have millions of records stored in a table in the program memory. This is definitely not realistic, thus we use neural network to estimate Q function and calculate the Q values.

1.2 b) Brief Description

The association of line numbers in code with lines in pseudocode is in the Appendix.

In the main function, first we initialize the CartPole-v0 environment, DQN agent and the arrays to collect test states for plotting Q values using random policy.

Then for every episode, we initialize the state first. And then get action for current state based on ϵ -greedy policy and go one step in environment, train the NN model, collect reward and store them and propagate the state, if it is not at the end of the episode. If it is at the end episode, we just update the target network and store the scores for plotting. If the check_solve flag is True, we calculate the mean of scores of last 100 episodes to check if it is larger than 195. If the average score is larger than 195, we stop the training process and mark as solved, otherwise we continue training.

The behavior of each function is:

__init__: the initialization of global variables.

build_model: use Keras to build the neural network model.

update_target_model: update target network model parameters, and make it the same with the model.

get_action: the implementation of ϵ -greedy policy, which is to determine how to select actions.

append_sample: save sample $\{s,a,r,s'\}$ to the replay memory.
train_model: implement the process of DQN.
plot_data: draw the image of average Q value and score.

1.3 c) Brief Explanation of Layout of NN

The layout of given Neural Network has two layers. The first layer has 16 neurons with nonlinear activation 'Relu', and the second layer has 2 neurons with a linear activation. Both layers use He Initialization to increase the stability and fasten the gradient descent algorithm.

1.4 d) *get-action* function

Please see the code in Appendix.

1.5 e) *train-model* function

Please see the code in Appendix.

1.6 f) Complexity of neural network

To compare the influence of the complexity of neural network, we first set the other hyper parameters to be fixed. Their values are:

discount factor = 0.95
learning rate = 0.005
epsilon = 0.02
batch size = 32
memory size = 1000
train start = 1000
target update frequency = 1

What is more, for each layer, the activation function is set as 'Relu' all the time. We change the structure of neural network for many times to evaluate the networks. Our result is shown from Figure 1 to Figure 5:

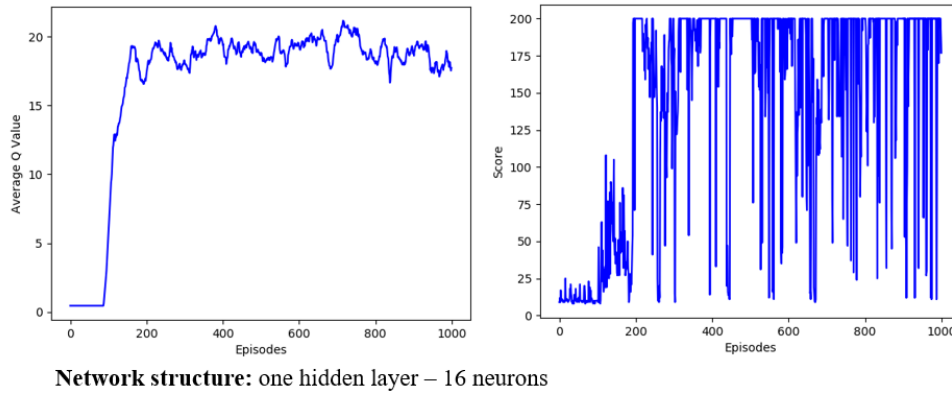




Figure 1: Network structure complexity case 1

As we can note, the agent learns faster with more nodes and more layers. From Figure 1 to Figure 5, the agent takes fewer steps to get high scores with more neurons since more layers and more nodes in the layer, we increase the complexity of the model, which has a better approximation of our problem. Thus we find the model with 32 neurons in the first layer and 32 neurons in the second layer the best and we will keep using this model in the following questions.




Network structure: one hidden layer – 32 neurons

Figure 2: Network structure complexity case 2



Network structure: two hidden layers – 16 neurons for the 1st layer, and 16 neurons for the 2nd layer

Figure 3: Network structure complexity case 3



Network structure: two hidden layers – 16 neurons for the 1st layer, and 32 neurons for the 2nd layer

Figure 4: Network structure complexity case 4


1.7 g) Effect of discount factor, learning rate, memory size

Discount factor

As we can observe from Figure 6 to Figure 8, the bigger discount factor is, the larger averaged Q value the agent can get, and the slower the average Q value will converge at the same time. Since small discount factor will force the agent focus on the current reward more, which is a disadvantage. So we decide to use discount factor 0.995 in the following questions.


Learning rate

As we can observe from Figure 9 to 11, a larger learning rate makes the agent learns faster but a




Network structure: two hidden layers – 32 neurons for the 1st layer, and 32 neurons for the 2nd layer

Figure 5: Network structure complexity case 5



Discount factor: 0.995

Figure 6: Discount factor case 1



Discount factor: 0.95

Figure 7: discount factor case 2

smaller learning rate makes the agent learns more stable. In order to have a relatively stable and fast learning process, we should choose the learning rate neither too large nor too small. Thus we decide to choose 0.005 as learning rate for following questions.

Memory size

As we can observe from Figure 12 to 14, when memory size is too small, namely smaller than 1000, the training process will not work. With larger memory size, since the agent can use more samples for training, it gives us better performance, but at the same time we observed that it took longer time to train the model. We will use memory size 10000 for the following questions.

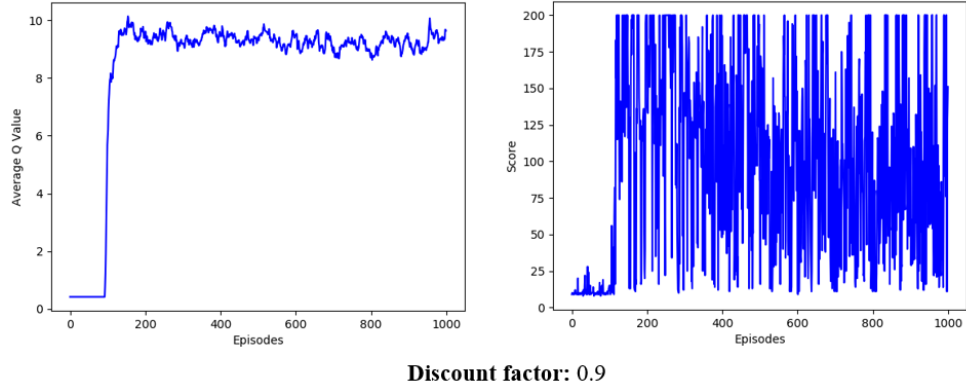


Figure 8: discount factor case 3




Figure 9: Learning rate case 1

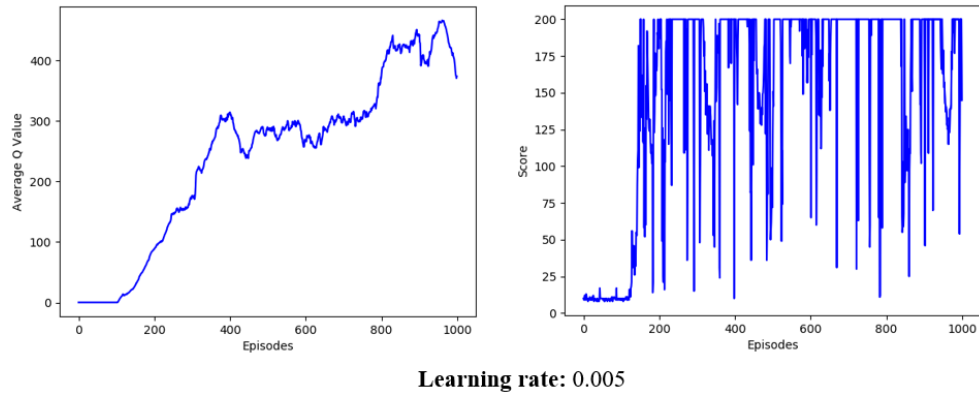



Figure 10: Learning rate case 2

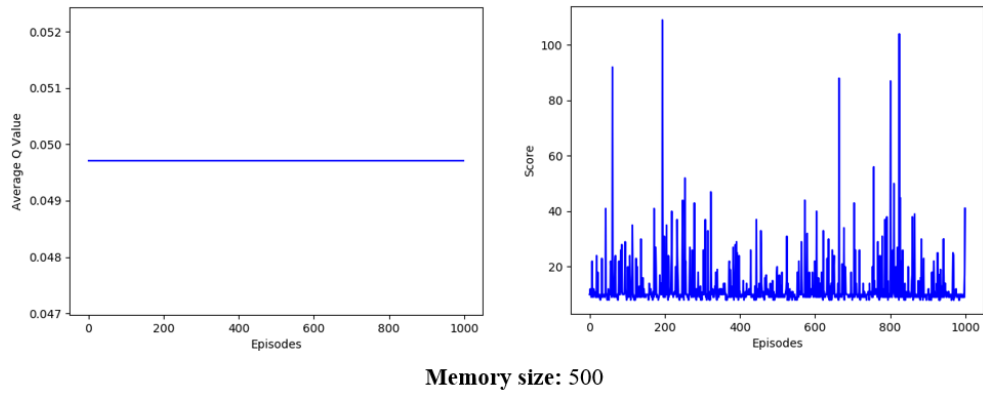
1.8 h) Effect of Target Update Frequency

As we can observe from Figure 15 to 18, target update frequency 1 gives us the best performance. The update frequency means how often we update the network, and a larger one will make the agent learn slower. From what we observe, frequency 1 gives the best performance. So we will use all improved hyper parameters to check if the problem solved.



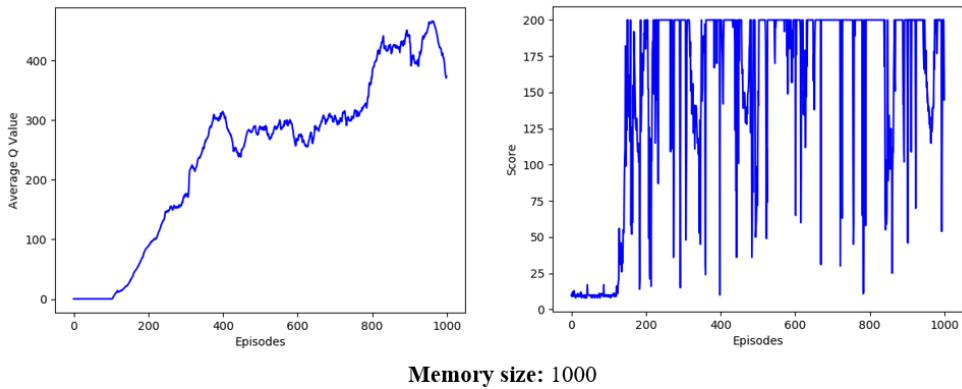
Learning rate: 0.002

Figure 11: Learning rate case 3




Memory size: 500

Figure 12: Memory size case 1




Memory size: 1000

Figure 13: Memory size case 2




Memory size: 10000

Figure 14: Memory size case 3



Target update frequency: 1

Figure 15: Target update frequency case 1



Target update frequency: 5

Figure 16: Target update frequency case 2

1.9 i) Performance of an Appropriate Set of Hyperparameters

Finally we choose:

discount factor = 0.995

learning rate = 0.005

epsilon = 0.02

batch size = 32

memory size = 10000

train start = 1000

target update frequency = 1

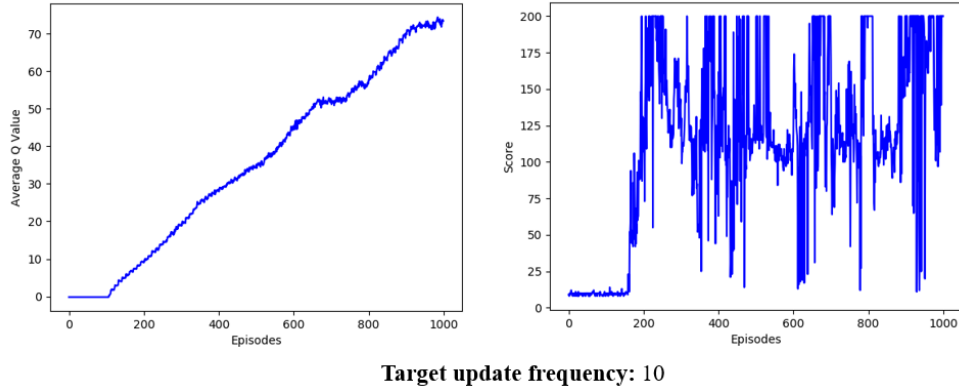


Figure 17: Target update frequency case 3

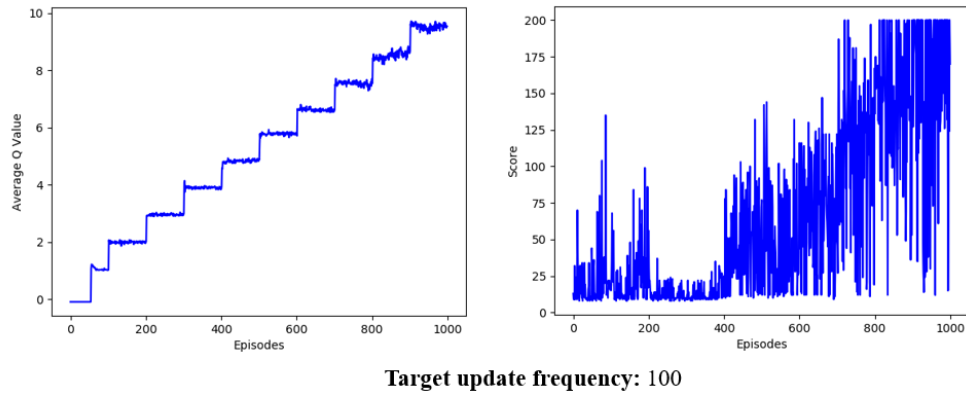


Figure 18: Target update frequency case 4

And observe that the problem solved after episode 724. The Q value and score plot can be found in Figure 19.




Figure 19: Performance of an Appropriate Set of Hyperparameters

2 Appendix

Line 203-215 in code is to initialize the first state and compute q-values, which corresponds to pseudocode line 1.

Line 216-219 in code is to do the iteration, which corresponds to pseudocode line 2.

Line 220-227 in code is to use ϵ -greedy policy to get action and update the next state, which corresponds to pseudocode line 3-5.

Line 229-231 in code is to train the agent, which corresponds to pseudocode line 6-12.

Line 233-242 in code is to update target network parameters 'frequency', which corresponds to pseudocode line 14.

Line 244-250 in code is to check if we can consider the problem is solved, by calculating the mean of scores of last 100 episodes, which corresponds to pseudocode line 15-17.

```
1 import sys
2 import gym
3 import pylab
4 import random
5 import numpy as np
6 from collections import deque
7 from keras.layers import Dense
8 from keras.optimizers import Adam
9 from keras.models import Sequential
10
11
12 EPISODES = 1000 #Maximum number of episodes
13
14 #DQN Agent for the Cartpole
15 #Q function approximation with NN, experience replay, and target network
16 class DQNAgent:
17     #Constructor for the agent (invoked when DQN is first called in main)
18     def __init__(self, state_size, action_size):
19         self.check_solve = True #If True, stop if you satisfy solution condition
20         self.render = False
21         #If you want to see Cartpole learning, then change to True
22
23         #Get size of state and action
24         self.state_size = state_size
25         self.action_size = action_size
26
27     #####
28     #####
29         #Set hyper parameters for the DQN. Do not adjust those labeled as Fixed.
30         self.discount_factor = 0.995
31         self.learning_rate = 0.002
32         self.epsilon = 0.02 #Fixed
33         self.batch_size = 32 #Fixed
34         self.memory_size = 10000
35         self.train_start = 1000 #Fixed
36         self.target_update_frequency = 1
37     #####
38     #####
39
40         #Number of test states for Q value plots
41         self.test_state_no = 10000
42
43         #Create memory buffer using deque
44         self.memory = deque(maxlen=self.memory_size)
45
46         #Create main network and target network (using build_model defined below)
47         self.model = self.build_model()
48         self.target_model = self.build_model()
49
50         #Initialize target network
```

```

51         self.update_target_model()
52
53     #Approximate Q function using Neural Network
54     #State is the input and the Q Values are the output.
55     #####
56     #####
57     #Edit the Neural Network model here
58     #Tip: Consult https://keras.io/getting-started/sequential-model-guide/
59     def build_model(self):
60         model = Sequential()
61         model.add(Dense(32, input_dim=self.state_size, activation='relu',
62                         kernel_initializer='he_uniform'))
63         model.add(Dense(32, activation='relu',
64                         kernel_initializer='he_uniform'))
65         model.add(Dense(self.action_size, activation='linear',
66                         kernel_initializer='he_uniform'))
67         # model.add(Dense(16, input_dim=self.state_size, activation='relu',
68         # kernel_initializer='he_uniform'))
69         # model.add(Dense(self.action_size, activation='linear',
70         # kernel_initializer='he_uniform'))
71         model.summary()
72         model.compile(loss='mse', optimizer=Adam(lr=self.learning_rate))
73         return model
74     #####
75     #####
76
77     #After some time interval update the target model to be same with model
78     def update_target_model(self):
79         self.target_model.set_weights(self.model.get_weights())
80
81     #Get action from model using epsilon-greedy policy
82     def get_action(self, state):
83     #####
84     #####
85         #Insert your e-greedy policy code here
86         #Tip 1: Use the random package to generate a random action.
87         #Tip 2: Use keras.model.predict() to compute Q-values from the state.
88         # action = random.randrange(self.action_size)
89         state_actions = self.model.predict(state) # All actions in this state
90         if (np.random.uniform() < self.epsilon) or (
91             state_actions.all() == 0):
92             # non-greedy or when this state has not been discovered
93             # if (np.random.uniform() <= self.epsilon):
94             action = random.randrange(self.action_size)
95         else:
96             action = state_actions.argmax() # greedy
97         return action
98     #####
99     #####
100     #Save sample <s,a,r,s'> to the replay memory
101     def append_sample(self, state, action, reward, next_state, done):
102         self.memory.append((state, action, reward, next_state, done))
103         #Add sample to the end of the list
104
105     #Sample <s,a,r,s'> from replay memory
106     def train_model(self):
107         if len(self.memory) < self.train_start: #Do not train if not enough memory
108             return
109         batch_size = min(self.batch_size, len(self.memory))

```

```

110         #Train on at most as many samples as you have in memory
111         mini_batch = random.sample(self.memory, batch_size)
112         #Uniformly sample the memory buffer
113         #Preallocate network and target network input matrices.
114         update_input = np.zeros((batch_size, self.state_size))
115         #batch_size by state_size two-dimensional array (not matrix!)
116         update_target = np.zeros((batch_size, self.state_size))
117         #Same as above, but used for the target network
118         action, reward, done = [], [], [] #Empty arrays that will grow dynamically
119
120         for i in range(self.batch_size):
121             update_input[i] = mini_batch[i][0]
122             #Allocate s(i) to the network input array from iteration i in the batch
123             action.append(mini_batch[i][1]) #Store a(i)
124             reward.append(mini_batch[i][2]) #Store r(i)
125             update_target[i] = mini_batch[i][3]
126             #Allocate s'(i) for the target network array
127             # from iteration i in the batch
128             done.append(mini_batch[i][4]) #Store done(i)
129
130         target = self.model.predict(update_input)
131         #Generate target values for training the inner loop network
132         # using the network model
133         target_val = self.target_model.predict(update_target)
134         #Generate the target values for training the outer loop target network
135
136         #Q Learning: get maximum Q value at s' from target network
137         #####
138         #####
139         #Insert your Q-learning code here
140         #Tip 1: Observe that the Q-values are stored in the variable target
141         #Tip 2: What is the Q-value of the action taken
142         # at the last state of the episode?
143         # for i in range(self.batch_size): #For every batch
144         #     target[i][action[i]] = random.randint(0,1)
145         for i in range(self.batch_size):
146             if done[i]:
147                 target[i][action[i]] = reward[i]
148             else:
149                 target[i][action[i]] = reward[i] + \
150                     self.discount_factor * np.max(target_val[i])
151         #####
152         #####
153
154         #Train the inner loop network
155         self.model.fit(update_input, target, batch_size=self.batch_size,
156                       epochs=1, verbose=0)
157         return
158     #Plots the score per episode as well as the maximum q value per episode,
159     # averaged over precollected states.
160     def plot_data(self, episodes, scores, max_q_mean):
161         pylab.figure(0)
162         pylab.plot(episodes, max_q_mean, 'b')
163         pylab.xlabel("Episodes")
164         pylab.ylabel("Average Q Value")
165         pylab.savefig("qvalues.png")
166
167         pylab.figure(1)
168         pylab.plot(episodes, scores, 'b')

```

```

169         pylab.xlabel("Episodes")
170         pylab.ylabel("Score")
171         pylab.savefig("scores.png")
172
173     if __name__ == "__main__":
174         #For CartPole-v0, maximum episode length is 200
175         env = gym.make('CartPole-v0')
176         #Generate Cartpole-v0 environment object from the gym library
177         #Get state and action sizes from the environment
178         state_size = env.observation_space.shape[0]
179         action_size = env.action_space.n
180
181         #Create agent, see the DQNAgent __init__ method for details
182         agent = DQNAgent(state_size, action_size)
183
184         #Collect test states for plotting Q values using uniform random policy
185         test_states = np.zeros((agent.test_state_no, state_size))
186         max_q = np.zeros((EPISODES, agent.test_state_no))
187         max_q_mean = np.zeros((EPISODES, 1))
188
189         done = True
190         for i in range(agent.test_state_no):
191             if done:
192                 done = False
193                 state = env.reset()
194                 state = np.reshape(state, [1, state_size])
195                 test_states[i] = state
196             else:
197                 action = random.randrange(action_size)
198                 next_state, reward, done, info = env.step(action)
199                 next_state = np.reshape(next_state, [1, state_size])
200                 test_states[i] = state
201                 state = next_state
202
203         scores, episodes = [], [] #Create dynamically growing score and episode counters
204         for e in range(EPISODES):
205             done = False
206             score = 0
207             state = env.reset() #Initialize/reset the environment
208             state = np.reshape(state, [1, state_size])
209             #Reshape state so that to a 1 by state_size two-dimensional array
210             # ie. [x_1, x_2] to [[x_1, x_2]]
211             #Compute Q values for plotting
212             tmp = agent.model.predict(test_states)
213             max_q[e][:] = np.max(tmp, axis=1)
214             max_q_mean[e] = np.mean(max_q[e][:])
215
216             while not done:
217                 if agent.render:
218                     env.render() #Show cartpole animation
219
220                 #Get action for the current state and go one step in environment
221                 action = agent.get_action(state)
222                 next_state, reward, done, info = env.step(action)
223                 next_state = np.reshape(next_state, [1, state_size])
224                 #Reshape next_state similarly to state
225
226                 #Save sample <s, a, r, s'> to the replay memory
227                 agent.append_sample(state, action, reward, next_state, done)

```

```

228         #Training step
229         agent.train_model()
230         score += reward #Store episodic reward
231         state = next_state #Propagate state
232
233     if done:
234         #At the end of every episode, update the target network
235         if e % agent.target_update_frequency == 0:
236             agent.update_target_model()
237         #Plot the play time for every episode
238         scores.append(score)
239         episodes.append(e)
240
241         print("episode:", e, "  score:", score, "  q_value:", max_q_mean[e],
242               "  memory length:", len(agent.memory))
243
244         # if the mean of scores of last 100 episodes is bigger than 195
245         # stop training
246         if agent.check_solve:
247             if np.mean(scores[-min(100, len(scores)):]) >= 195:
248                 print("solved after", e-100, "episodes")
249                 agent.plot_data(episodes, scores, max_q_mean[:e+1])
250                 sys.exit()
251
252     pylab.figure()
253     pylab.plot(episodes[-min(100, len(scores)):],
254               scores[-min(100, len(scores)):], 'b')
255     pylab.xlabel("Episodes")
256     pylab.ylabel("Average Score in last 100 episode")
257     pylab.savefig("AverageScoresLast100Episode.png")
258     print( np.mean(scores[-min(100, len(scores)):]))
259     agent.plot_data(episodes, scores, max_q_mean)

```
